

Recitation 3: Geometric Derivation of SVMs

Intro Question

1. You have been given a data set (x_i, y_i) for $i = 1, \dots, n$ where $x_i \in \mathbb{R}^d$ and $y_i \in \{-1, 1\}$. Assume $w \in \mathbb{R}^d$ and $a \in \mathbb{R}$.
 - (a) Suppose $y_i(w^T x_i + a) > 0$ for all i . Use a picture to explain what this means when $d = 2$.
 - (b) Fix $M > 0$. Suppose $y_i(w^T x_i + a) \geq M$ for all i . Use a picture to explain what this means when $d = 2$.




Figure 1: Data set with $x_i \in \mathbb{R}^2$ and $y_i \in \{+1, -1\}$

Support Vector Machines

Review of Geometry

If $v, w \in \mathbb{R}^d$ then the component (also called scalar projection) of v in the direction w is given by the scalar $\frac{w^T v}{\|w\|_2}$. This can also be thought of as the signed length of v when orthogonally projected onto the line through the vector w .




Figure 2: Component of v_1, v_2 in the direction w

Assuming $w \neq 0$ we can use this to interpret the set

$$S = \{x \in \mathbb{R}^d \mid w^T x = b\}.$$

Noting that $w^T x = b \iff \frac{w^T x}{\|w\|_2} = \frac{b}{\|w\|_2}$ we see that S contains all vectors whose component in the direction w is $\frac{b}{\|w\|_2}$. Using linear algebra we can see this is the hyperplane orthogonal to the vector w that passes through the point $\frac{bw}{\|w\|_2^2}$. Note also that there are infinitely many pairs (w, b) that give the same hyperplane. If $c \neq 0$ then

$$\{x \in \mathbb{R}^d \mid w^T x = b\} \quad \text{and} \quad \{x \in \mathbb{R}^d \mid (cw)^T x = (cb)\}$$

result in the same hyperplanes.




Figure 3: Level Surfaces of $f(v) = w^T v$ with $\|w\|_2 = 1$

Given a hyperplane $\{v \mid w^T v = b\}$, we can distinguish points $x \in \mathbb{R}^d$ depending on whether $w^T x - b$ is zero, positive, or negative, or in other words, whether x is on the hyperplane, on the side w is pointing at, or on the side $-w$ is pointing at.




Figure 4: Sides of the Hyperplane $w^T v = 15$

If we have a vector $x \in \mathbb{R}^d$ and a hyperplane $H = \{v \mid w^T v = b\}$ we can measure the distance from x to H by

$$d(x, H) = \left| \frac{w^T x - b}{\|w\|_2} \right|.$$

Without the absolute values we get the *signed distance*: a positive distance if $w^T x > b$ and a negative distance if $w^T x < b$. To see why this formula is correct, note that we are computing

$$\frac{w^T x}{\|w\|_2} - \frac{w^T v}{\|w\|_2},$$

where v is any vector in the hyperplane $\{v \mid w^T v = b\}$. This is the difference between their components in the direction w .




Figure 5: Signed Distance from x_1, x_2 to Hyperplane $w^T v = 20$

Hard Margin SVM

Returning to the initial question, suppose we have the data set (x_i, y_i) for $i = 1, \dots, n$ where $x_i \in \mathbb{R}^d$ and $y_i \in \{-1, 1\}$.

Definition 1 (Linearly Separable). We say (x_i, y_i) for $i = 1, \dots, n$ are *linearly separable* if there is a $w \in \mathbb{R}^d$ and $a \in \mathbb{R}$ such that $y_i(w^T x_i + a) > 0$ for all i . The set $\{v \in \mathbb{R}^d \mid w^T v + a = 0\}$ is called a *separating hyperplane*.

Let's examine what this definition says. If $y_i = +1$ then we require that $w^T x_i > -a$ and if $y_i = -1$ we require that $w^T x_i < -a$. Thus linearly separable means that we can separate all of the $+1$'s from the -1 's using the hyperplane $\{v \mid w^T v = -a\}$. For the rest of this section, we assume our data is linearly separable.




Figure 6: Linearly Separable Data

If we can find the w, a corresponding to a hyperplane that separates the data, we then have a decision function for classifying elements of \mathcal{X} : $f(x) = \text{sgn}(w^T x + a)$. Before we look for such a hyperplane, we must address another issue. If the data is linearly separable, then there are infinitely many choices of separating hyperplanes.




Figure 7: Many Separating Hyperplanes Exist

We will choose the hyperplane that maximizes a quantity called the *geometric margin*.

Definition 2 (Geometric Margin). Let H be a hyperplane that separates the data (x_i, y_i)

for $i = 1, \dots, n$. The geometric margin of this hyperplane is

$$\min_i d(x_i, H),$$

the distance from the hyperplane to the closest data point.

Fix $w \in \mathbb{R}^d$ and $a \in \mathbb{R}$ with $y_i(w^T x_i + a) > 0$ for all i . Then we saw earlier that

$$d(x_i, H) = \left| \frac{w^T x_i + a}{\|w\|_2} \right| = \frac{y_i(w^T x_i + a)}{\|w\|_2}.$$

This gives us the following optimization problem:

$$\text{maximize}_{w,a} \quad \min_i \frac{y_i(w^T x_i + a)}{\|w\|_2}.$$

We can rewrite this in a more standard form:

$$\begin{aligned} & \text{maximize}_{w,a,M} \quad M \\ & \text{subject to} \quad \frac{y_i(w^T x_i + a)}{\|w\|_2} \geq M \quad \text{for all } i. \end{aligned}$$




Figure 8: Maximum Margin Separating Hyperplane

Note above how the geometric margin is achieved on both sides of the optimal hyperplane. This must be the case, as otherwise we could slightly move the hyperplane and obtain a better solution. The expression $y_i(w^T x_i + a)/\|w\|_2$ allows us to choose any positive value for $\|w\|_2$ by changing a accordingly (e.g., we can replace $w \rightarrow 2w$ and $a \rightarrow 2a$ and get the same value for all (x_i, y_i)). Thus we can fix $\|w\|_2 = 1/M$ and obtain

$$\begin{aligned} & \text{maximize}_{w,a} && 1/\|w\|_2 \\ & \text{subject to} && y_i(w^T x_i + a) \geq 1 \quad \text{for all } i. \end{aligned}$$

To find the optimal w, a we can instead solve the minimization problem

$$\begin{aligned} & \text{minimize}_{w,a} && \|w\|_2^2 \\ & \text{subject to} && y_i(w^T x_i + a) \geq 1 \quad \text{for all } i. \end{aligned}$$

This is a quadratic program that can be solved by standard packages.

Here the geometric margin is $\frac{1}{\|w\|_2}$ which is also the minimum of $\frac{y_i(w^T x_i + a)}{\|w\|_2}$ over the training data. The concept of margin used in class (also called *functional margin*) is slightly different, but clearly related. It is the value $y_i(w^T x_i + a)$ denoting the score we give to a given training example.

Soft Margin SVM

The methods developed thus far require linearly separable data. To remove this restriction, we will allow vectors to violate the geometric margin requirements, but at a penalty. More precisely, we replace our previous SVM formulation

$$\begin{aligned} & \text{minimize}_{w,a} && \|w\|_2^2 \\ & \text{subject to} && y_i(w^T x_i + a) \geq 1 \quad \text{for all } i \end{aligned}$$

with

$$\begin{aligned} & \text{minimize}_{w,a,\xi} && \|w\|_2^2 + \frac{C}{n} \sum_{i=1}^n \xi_i \\ & \text{subject to} && y_i(w^T x_i + a) \geq 1 - \xi_i \quad \text{for all } i \\ & && \xi_i \geq 0 \quad \text{for all } i. \end{aligned}$$

This is the standard formulation of a support vector machine. When $\xi_i > 0$ the corresponding x_i violates the geometric margin condition. Each ξ_i is called a slack variable. The constant C controls how much we penalize violations. Rewriting the condition as

$$\frac{y_i(w^T x_i + a)}{\|w\|_2} \geq \frac{1 - \xi_i}{\|w\|_2}$$

shows that ξ_i measures the size of the violation in multiples of the geometric margin. For example, $\xi_i = 1$ means x_i lies on the decision hyperplane $w^T v + a = 0$, and $\xi_i = 3$ means x_i lies 2 margin widths past the decision hyperplane.




Figure 9: Soft Margin SVM (unlabeled points have $\xi_i = 0$)

Recall from the treatment in class, that the minimizer w will be a linear combination of some of the x_i , called support vectors. More precisely, the support vectors will be some subset of the x_i that either lie on the margin boundary ($y_i(w^T x_i + a) = 1$) or violate the margin boundary ($y_i(w^T x_i + a) < 1$, $\xi_i > 0$).

Regularization Interpretation

Consider the following two questions:

1. If your data is linearly separable, which SVM (hard margin or soft margin) would you use?
2. Explain geometrically what the following optimization problem computes:

$$\begin{aligned}
 & \text{minimize}_{w,a,\xi} && \frac{1}{n} \sum_{i=1}^n \xi_i \\
 & \text{subject to} && y_i(w^T x_i + a) \geq 1 - \xi_i \quad \text{for all } i \\
 & && \|w\|_2^2 \leq r^2 \\
 & && \xi_i \geq 0 \quad \text{for all } i.
 \end{aligned}$$

By dividing the soft margin objective by C and writing $\lambda = 1/C$ we obtain the equivalent minimization problem

$$\begin{aligned} & \text{minimize}_{w,a,\xi} && \lambda \|w\|_2^2 + \frac{1}{n} \sum_{i=1}^n \xi_i \\ & \text{subject to} && y_i(w^T x_i + a) \geq 1 - \xi_i \quad \text{for all } i \\ & && \xi_i \geq 0 \quad \text{for all } i. \end{aligned}$$

This has the form of a regularized objective where the average slack is the loss and $\lambda \|w\|_2^2$ is the L_2 regularization. By choosing λ we determine the trade-off between minimizing the slack of the violations, while keeping the other points at a reasonable margin from the decision boundary. As with linear regression, there is an equivalent Ivanov formulation:

$$\begin{aligned} & \text{minimize}_{w,a,\xi} && \frac{1}{n} \sum_{i=1}^n \xi_i \\ & \text{subject to} && y_i(w^T x_i + a) \geq 1 - \xi_i \quad \text{for all } i \\ & && \|w\|_2^2 \leq r^2 \\ & && \xi_i \geq 0 \quad \text{for all } i. \end{aligned}$$

Recall that the geometric margin is $1/\|w\|_2$. Thus the Ivanov regularized problem is to minimize the average slack, but only among classifiers that have a margin of at least $1/r$.




Figure 10: Optimize Over Cases Where Margin Is At Least $1/r$

To see the value of regularization, consider the following examples.




Figure 11: Overfitting: Tight Margin With No Misclassifications




Figure 12: Training Error But Large Margin

Although the first figure above has no misclassifications on the training set, it has an incredibly small geometric margin. As a result, we may be willing to suffer a single training mistake in return for a large buffer region for most training examples. Note that the data here was linearly separable, but we still may prefer the soft margin SVM. By using cross-validation we can use data to find the correct tradeoff.